

# Revision


Lecture 12

February 14, 2024


# Broadcast Algorithms in MPICH

- Short messages
  - `< MPIR_CVAR_BCAST_SHORT_MSG_SIZE`
  - Binomial
- Medium messages
  - Scatter + Allgather (Recursive doubling)
- Large messages
  - `> MPIR_CVAR_BCAST_LONG_MSG_SIZE`
  - Scatter + Allgather (Ring)

# Old vs. New MPI\_Bcast




# Reduce on 64 nodes



# Allgather – Ring Algorithm

- Every process sends to and receives from everyone else
- Assume  $p$  processes and total  $n$  bytes
- Every process sends and receives  $n/p$  bytes
- Time
  - $(p - 1) * (L + n/p * (1/B))$
- How can we improve?



# Non-blocking Point-to-Point

- MPI\_Isend (buf, count, datatype, dest, tag, comm, request)
- MPI\_Irecv (buf, count, datatype, source, tag, comm, request)
- MPI\_Wait (request, status)
- MPI\_Waitall (count, request, status)

# Many-to-one Non-blocking P2P

```
// send from all ranks to the last rank
start_time = MPI_Wtime ();
if (myrank < size-1)
{
    MPI_Send(arr, BUFSIZE, MPI_INT, size-1, 99, MPI_COMM_WORLD);
}
else
{
    int count, recvarr[size][BUFSIZE];
    for (int i=0; i<size-1; i++)
        MPI_Irecv(recvarr[i], BUFSIZE, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &request[i]);
    MPI_Waitall (size-1, request, status);
}
time = MPI_Wtime () - start_time;

MPI_Reduce (&time, &max_time, 1, MPI_DOUBLE, MPI_MAX, size-1, MPI_COMM_WORLD);
if (myrank == size-1) printf ("Max time = %lf\n", max_time);
```

# Non-blocking Performance

- Standard does not require overlapping communication and computation
- Implementation **may** use a thread to move data in parallel
- Implementation **can delay** the initiation of data transfer until “Wait”
- MPI\_Test – non-blocking, tests completion, starts progress
- MPIR\_CVAR\_ASYNC\_PROGRESS (MPICH)




# Non-blocking Point-to-Point Safety

- MPI\_Isend (buf, count, datatype, dest, tag, comm, request)
- MPI\_Irecv (buf, count, datatype, source, tag, comm, request)
- MPI\_Wait (request, status)




# Mesh Interconnect



- Diameter  $2(\sqrt{p} - 1)$
- Bisection width  $\sqrt{p}$
- Cost  $2(p - \sqrt{p})$

# Torus Interconnect



- Diameter  $2(\sqrt{p}/2)$
- Bisection width  $2\sqrt{p}$
- Cost  $2p$

# Parallelization

# Parallelization Steps

## 1. *Decomposition* of computation into tasks

- Identifying portions of the work that can be performed concurrently

## 2. *Assignment* of tasks to processes

- Assigning concurrent pieces of work onto multiple processes running in parallel

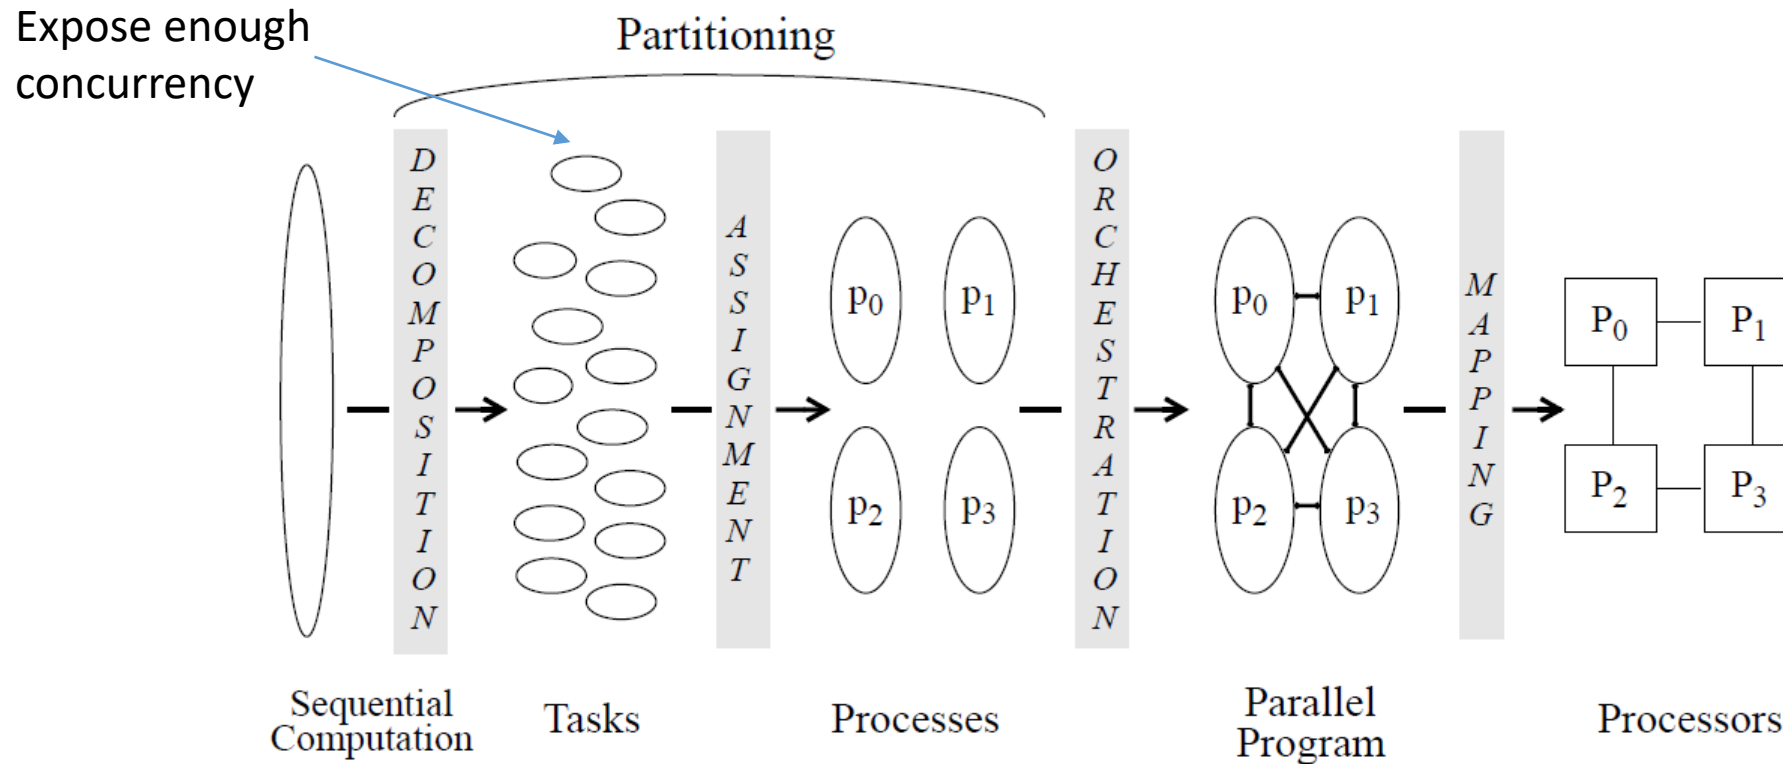
## 3. *Orchestration* of data access, communication and synchronization among processes

- Distributing the data associated with the program
- Managing access to data shared by multiple processes
- Synchronizing at various stages of the parallel program execution

## 4. *Mapping* of processes to processors

- Placement of processes in the physical processor topology

# Illustration of Parallelization Steps




Source: Culler et al. book

# Performance Goals

- Expose concurrency
- Reduce inter-process communications
- Load-balance
- Reduce synchronization
- Reduce idling
- Reduce management overhead
- Preserve data locality
- Exploit network topology

# Matrix Vector Multiplication – Decomposition



$P = 3 ?$

Decomposition


Identifying portions of the  
work that can be performed  
concurrently

Assignment



# Matrix Vector Multiplication – Orchestration

$P = 3$



Decomposition

Assignment

Orchestration

- Allgather/Bcast
- Scatter
- Gather

- Initial communication
  - Distribute (read by process 0) or parallel reads
- Final communication

# Distribute using Bcast vs. Allgather

```
MPI_Init( &argc, &argv );
MPI_Comm_rank( MPI_COMM_WORLD, &myrank );
MPI_Comm_size( MPI_COMM_WORLD, &commsize );

time = MPI_Wtime();
MPI_Bcast (buf, N, MPI_FLOAT, 0, MPI_COMM_WORLD);
etime = MPI_Wtime() - time;
MPI_Reduce (&etime, &maxtime, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
if (!myrank) printf ("Time to bcast: %11.3lf\n", maxtime);

int count = N/commsize;
time = MPI_Wtime();
MPI_Allgather (buf, count, MPI_FLOAT, recvbuf, count, MPI_FLOAT, MPI_COMM_WORLD);
etime = MPI_Wtime() - time;
MPI_Reduce (&etime, &maxtime, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
if (!myrank) printf ("Time to allgather: %7.3lf\n", maxtime);
```

```
class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 10000; echo ; done
Time to bcast:      0.014
Time to allgather:  0.021

Time to bcast:      0.018
Time to allgather:  0.009

Time to bcast:      0.012
Time to allgather:  0.007

class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 100000; echo ; done
Time to bcast:      0.034
Time to allgather:  0.011

Time to bcast:      0.027
Time to allgather:  0.023

Time to bcast:      0.026
Time to allgather:  0.011

class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 1000000; echo ; done
Time to bcast:      0.187
Time to allgather:  0.347

Time to bcast:      0.176
Time to allgather:  0.111

Time to bcast:      0.155
Time to allgather:  0.112
```

```
class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 1000000; echo ; done
Time to bcast:      0.187
Time to allgather:  0.347

Time to bcast:      0.176
Time to allgather:  0.111

Time to bcast:      0.155
Time to allgather:  0.112

class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 10000000; echo ; done
Time to bcast:      1.421
Time to allgather:  1.121

Time to bcast:      1.618
Time to allgather:  1.282


Time to bcast:      1.674
Time to allgather:  1.583

class for i in `seq 1 3` ; do mpirun -np 10 -hosts csews2,csews5,csews20 ./bcast-allgather 100000000; echo ; done
Time to bcast:      18.061
Time to allgather:  15.616

Time to bcast:      23.447
Time to allgather:  17.005

Time to bcast:      16.875
Time to allgather:  11.085
```

# Matrix Vector Multiplication – Column-wise Decomposition




Decomposition  
Assignment  
Orchestration

- Reduce

Row-wise vs. column-wise partitioning?

# 1D Domain Decomposition



Nearest neighbor communications

```
2 sends()  
2 recvs()
```

N grid points

P processes

$N/P$  points per process

#Communications?


2

#Computations?

$N/P$

Communication to computation ratio =  $2P/N$

# 1D Domain Decomposition



$N$  grid points  
 $P$  processes  
 $N/P$  points per process

**#Communications?**


$2\sqrt{N}$  (assuming square grid)

**#Computations?**

$N/P$  (assuming square grid)

Communication to computation ratio=?

# 2D Domain decomposition




$N$  grid points ( $\sqrt{N} \times \sqrt{N}$  grid)  
 $P$  processes ( $\sqrt{P} \times \sqrt{P}$  grid)  
 $N/P$  points per process

- + Several parallel communications
- + Lower communication volume/process



# 2D Domain decomposition



2 Sends()  
2 Recvs()

*#Communications?*


$2\sqrt{N}/\sqrt{P}$  (assuming square grid)

*#Computations?*

$N/P$  (assuming square grid)


*Communication to computation ratio=?*

# Stencils



Five-point stencil

# 2D Domain decomposition



4 Sends()  
4 Recvs()

N grid points ( $\sqrt{N} \times \sqrt{N}$  grid)  
P processes ( $\sqrt{P} \times \sqrt{P}$  grid)  
N/P points per process

**#Communications?**


$4\sqrt{N}/\sqrt{P}$  (assuming square grid)

**#Computations?**

N/P (assuming square grid)


Communication to computation ratio=?

# Send / Recv




MPI\_Send

MPI\_Recv




# Send / Recv



MPI\_Pack (buf)  
MPI\_Send (buf)

MPI\_Recv (buf)  
MPI\_Unpack (buf)



# MPI\_Pack

```
MPI_Init(&argc, &argv);

MPI_Comm_rank(MPI_COMM_WORLD, &myrank) ;
MPI_Comm_size(MPI_COMM_WORLD, &size);

// initialize data
for (int i=0; i<M; i++)
    for (int j=0; j<N; j++)
        array2D[i][j] = myrank+i+j;

sTime = MPI_Wtime();
if (myrank == 0) {
    // pack the last element of every row (N ints)
    for (int j=0; j<N; j++) {
        MPI_Pack (&array2D[j][M-1], 1, MPI_INT, buffer, 400, &position, MPI_COMM_WORLD);
        printf ("packed %d %d\n", j, position);
    }
    MPI_Send (buffer, position, MPI_PACKED, 1, 1, MPI_COMM_WORLD);
}
else {
    // receive N ints
    if (myrank == 1)
        MPI_Recv (buffer, count, MPI_INT, 0, 1, MPI_COMM_WORLD, &status);
    // verify
    MPI_Get_count (&status, MPI_INT, &count);
}
eTime = MPI_Wtime();
time = eTime - sTime;

printf ("%lf\n", time);
```

int MPI\_Pack (const void  
\*inbuf, int incount,  
MPI\_Datatype datatype,  
void \*outbuf, int outsize,  
int \*position, MPI\_Comm  
comm)