

# Paper Discussion

# Topology and Interference

Lecture 22

Apr 10, 2024

Network Performance Aware MPI Collective  
Communication Operations in the Cloud  
Gong et al., IEEE TPDS 2015

# Differences/Challenges

- Topology information is not available
- Certain privileges are not available
- Topology is not static
- Network asymmetry
  - Performance varies for different VM pairs
- Schedule communication for machine pairs with low network performance in an optimal manner
- How do we infer/discover the topology

# Goals

- Model to capture the network performance among different machine pairs in the cloud environment
- Network performance aware algorithms for collective operations with different message sizes
  - Broadcast
  - Reduce


# Topology Discovery

- Communication time between  $N/2$  pairs are measured (repeated)
- Latency matrix  $[N \times N]$ 
  - $L(i, j)$  – average latency from node  $i$  to node  $j$
  - $L(i, i) = 0$
- Bandwidth matrix  $[N \times N]$ 
  - $B(i, j)$  – average bandwidth from node  $i$  to node  $j$
  - $B(i, i) = ?$
- Network performance matrix for a message size  $m$   $[N \times N]$ 
  - $M(i, j) = L(i, j) + m/B(i, j)$

# Algorithm

- Network performance among a set of VMs
  - Latency matrix
  - Bandwidth matrix
  - Network performance matrix (Hockney model)
- Group the machines
  - Form a number of non-overlapping groups
  - Each group has machines that have similar network performance (use a distance parameter and threshold)
  - Build hierarchy among the groups by increasing the closeness parameter
  - Adjust span (difference in #machines between largest and smallest group)

# Network-aware Hierarchy




# Network Performance Aware Reduce

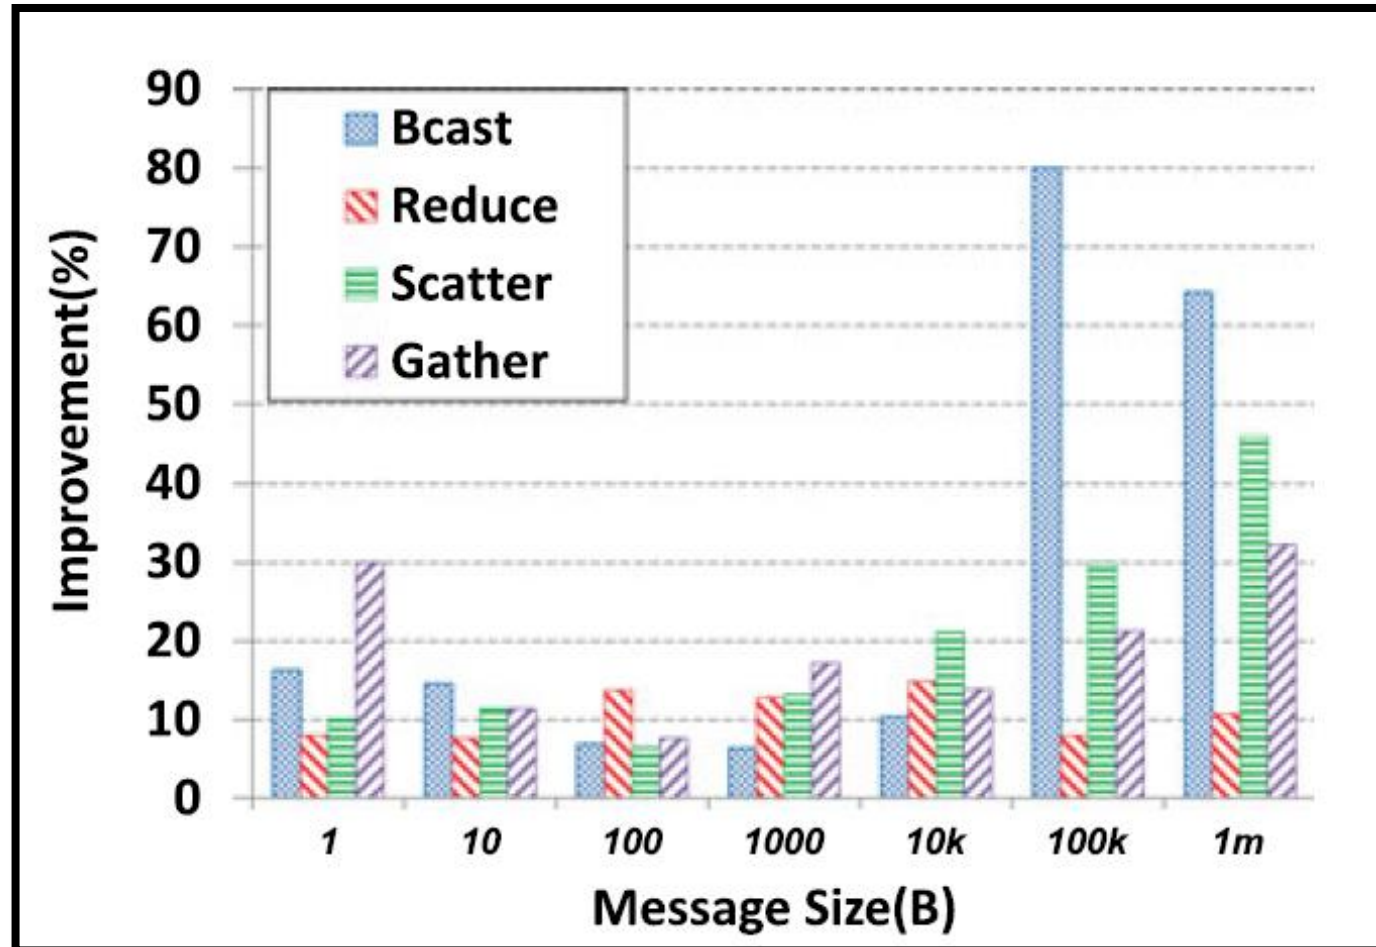
- Build network performance hierarchy based on L, B, m (message size)
- Select leader process in each group
- Create sub-communicators for each group
- Each group performs reduce independently
- Hierarchically reduce



# Network-aware Reduce



# Results



# Analyzing Network Health and Congestion in Dragonfly-based Supercomputers Bhatele et al., IPDPS 2016

# Overview

- Job-interference is an important factor for communication-intensive applications
  - Network congestion
- Understanding causes of congestion
  - Routing algorithms
  - Job placement policies
  - Network topology
- 44 workloads, on up to 131K cores, 588 simulations exploring different causes

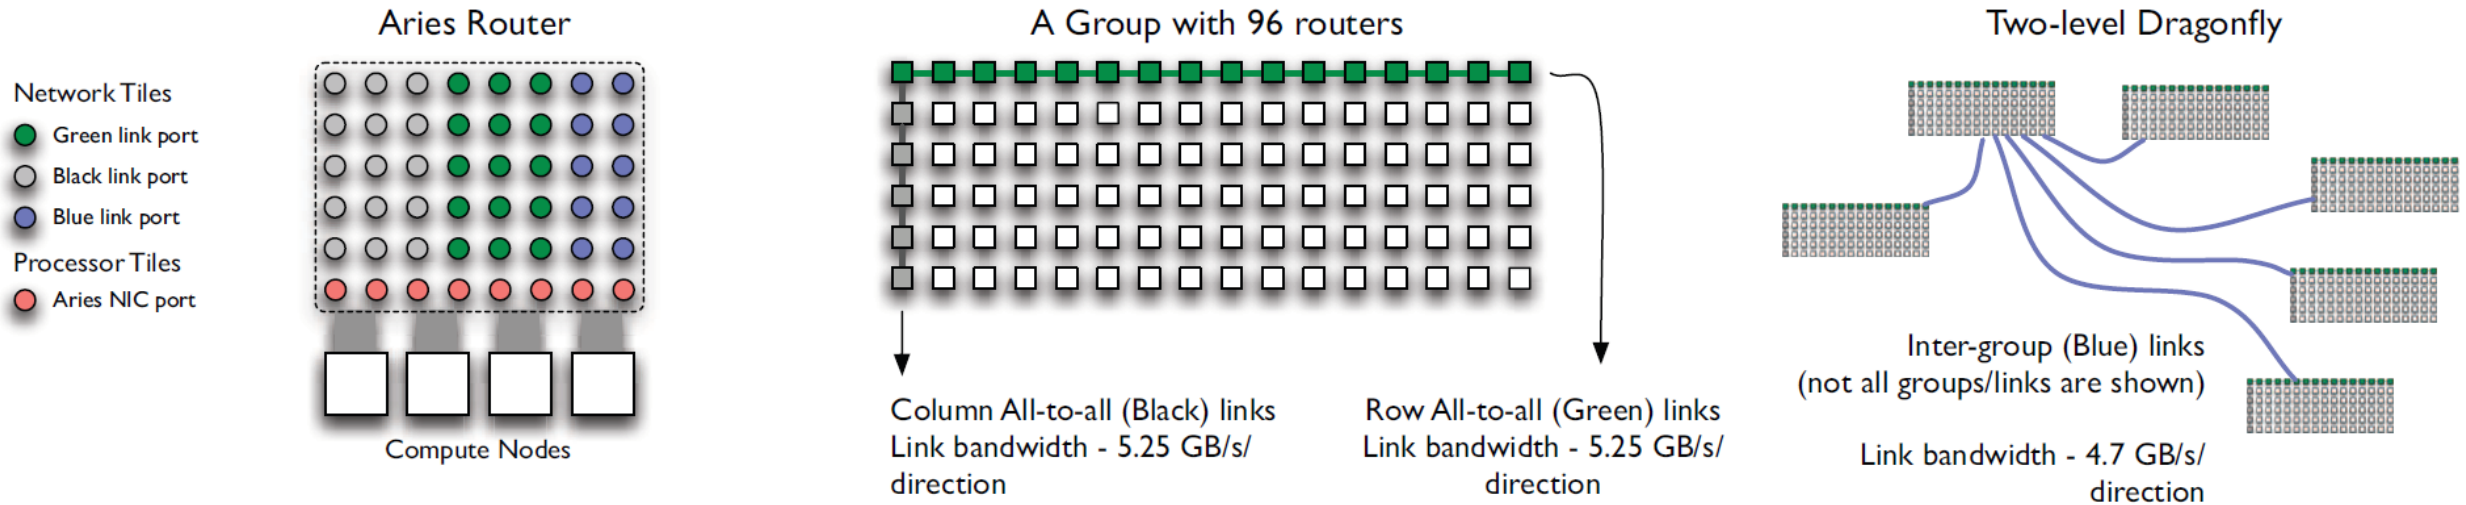
# Network Simulation

- Parallel network simulator
- Simulate parallel workloads and record network traffic (Damselfly: <https://github.com/LLNL/damselfly>)
- Input
  - Router connectivity graph
  - Application communication graph
  - Routing policy
- Output
  - Network link traffic (counters) for multi-job simulations
  - Identify job interference
- Enables
  - Exploring job placements, routing, network configurations

# Communication Patterns


- 2D stencil (4 neighbours, 64 KB)
- 4D stencil (8 neighbours, 4 MB)
- Many-to-many (all to all in groups of 16 or 32, 32 KB)
- Spread (Random 6 to 27 neighbours, 512 KB)
- Unstructured mesh (Select 6 to 27 neighbours, 512 KB)

# Dragonfly Network Topology



# Link Prediction (Isolated)

Average and Maximum Traffic for Individual Jobs (Green, Black and Blue links)

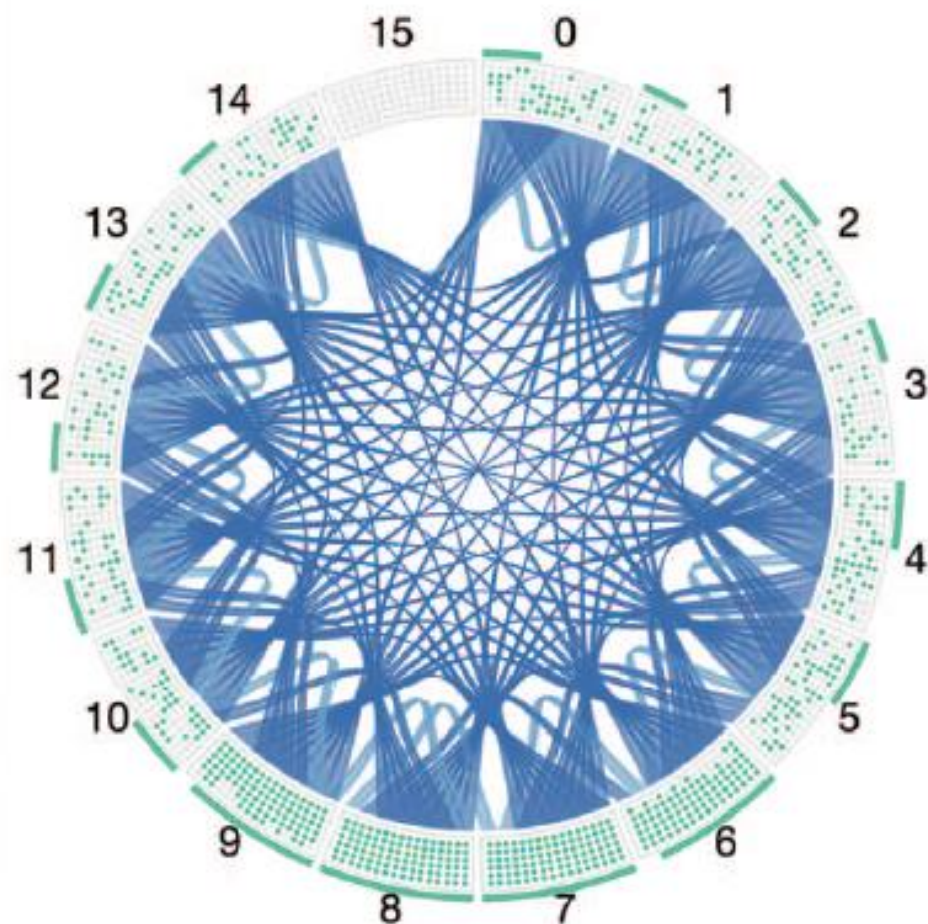
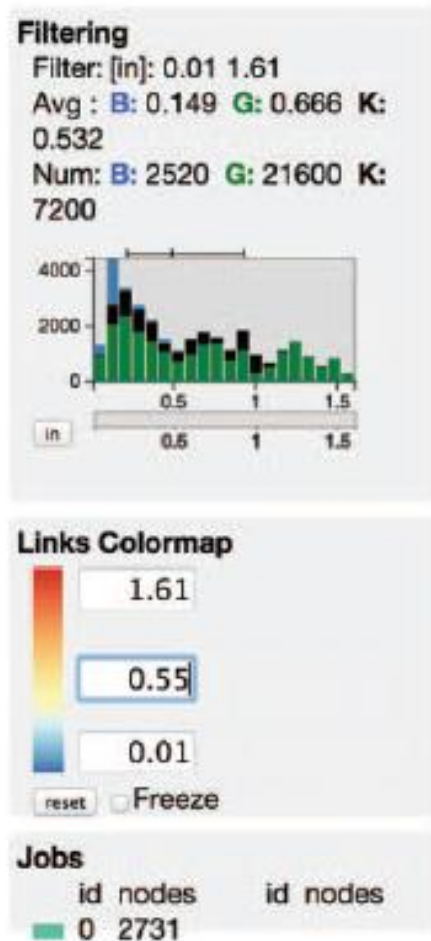


## Observations

- Link traffic increases with job size
- Some create significantly high traffic
- Maximum traffic on ? links → hot-spots





# Isolated Network Traffic (64K cores)

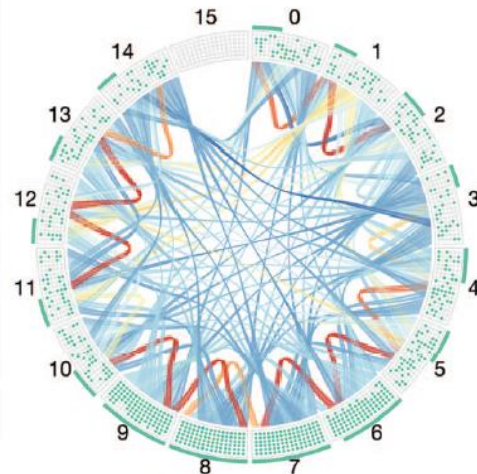



(a) 2D Stencil



# Isolated Network Traffic (64K cores)



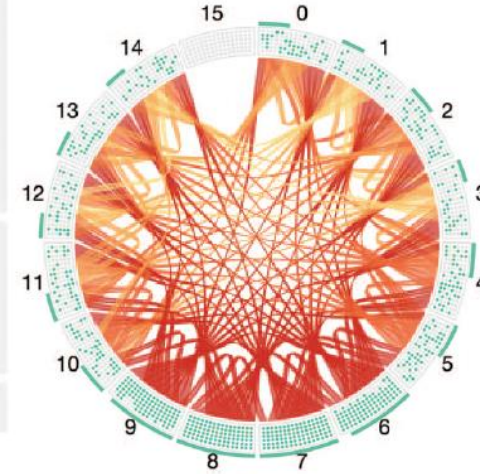

(a) 2D Stencil



(c) Many-to-many



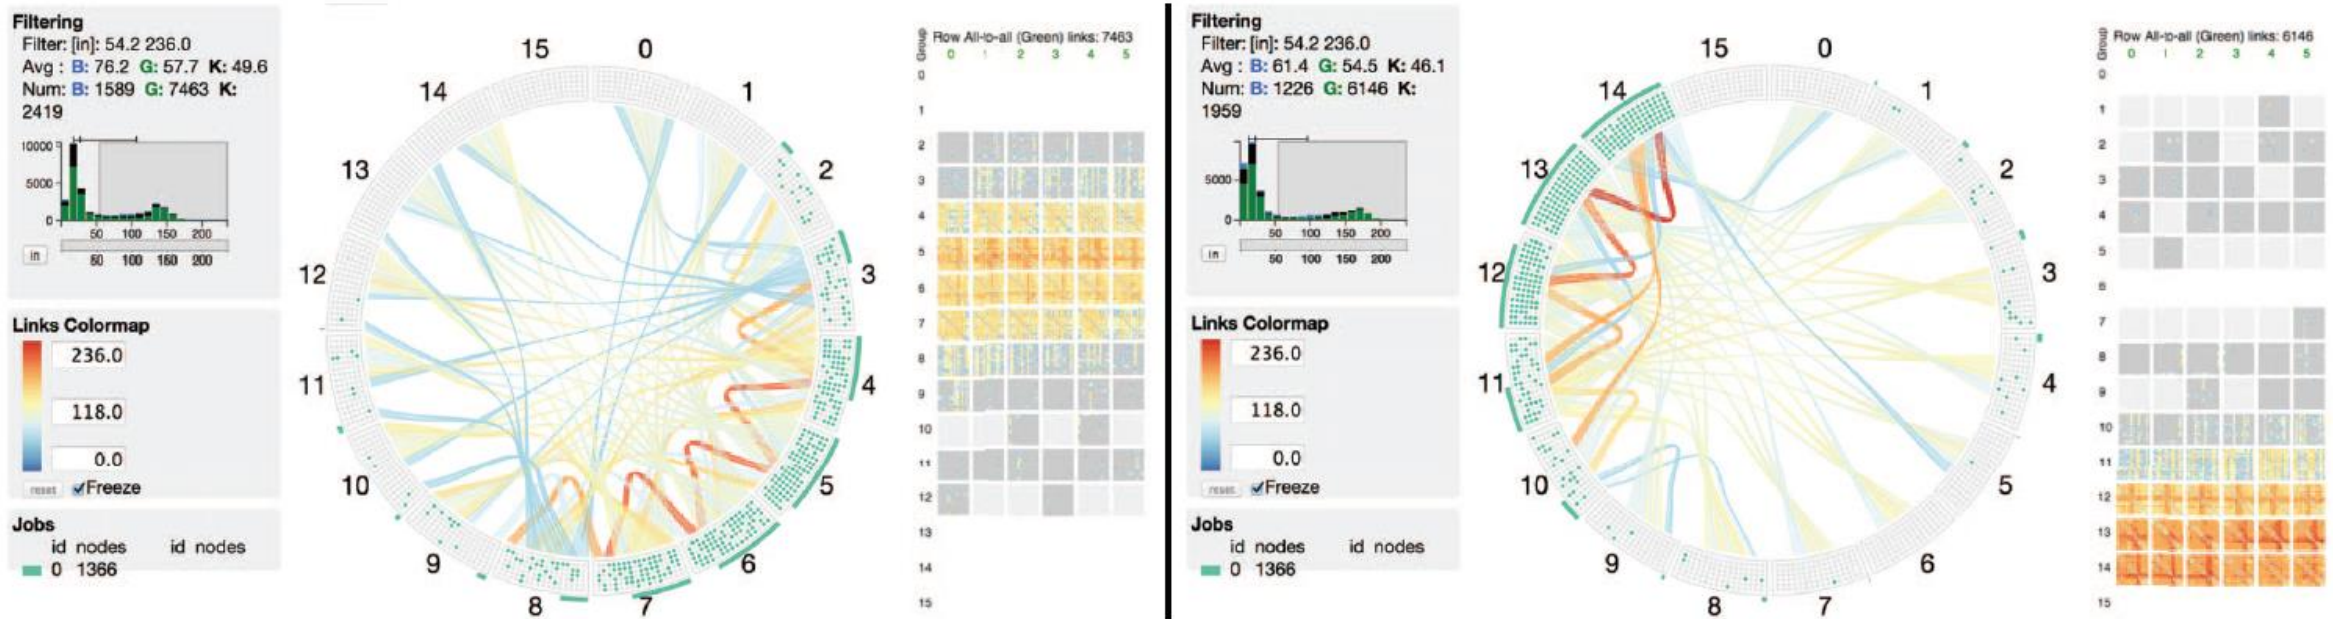
(b) 4D Stencil



(d) Spread



# Job Placement Effect



Scattered vs. Compact placement of 4D stencil on 32K cores

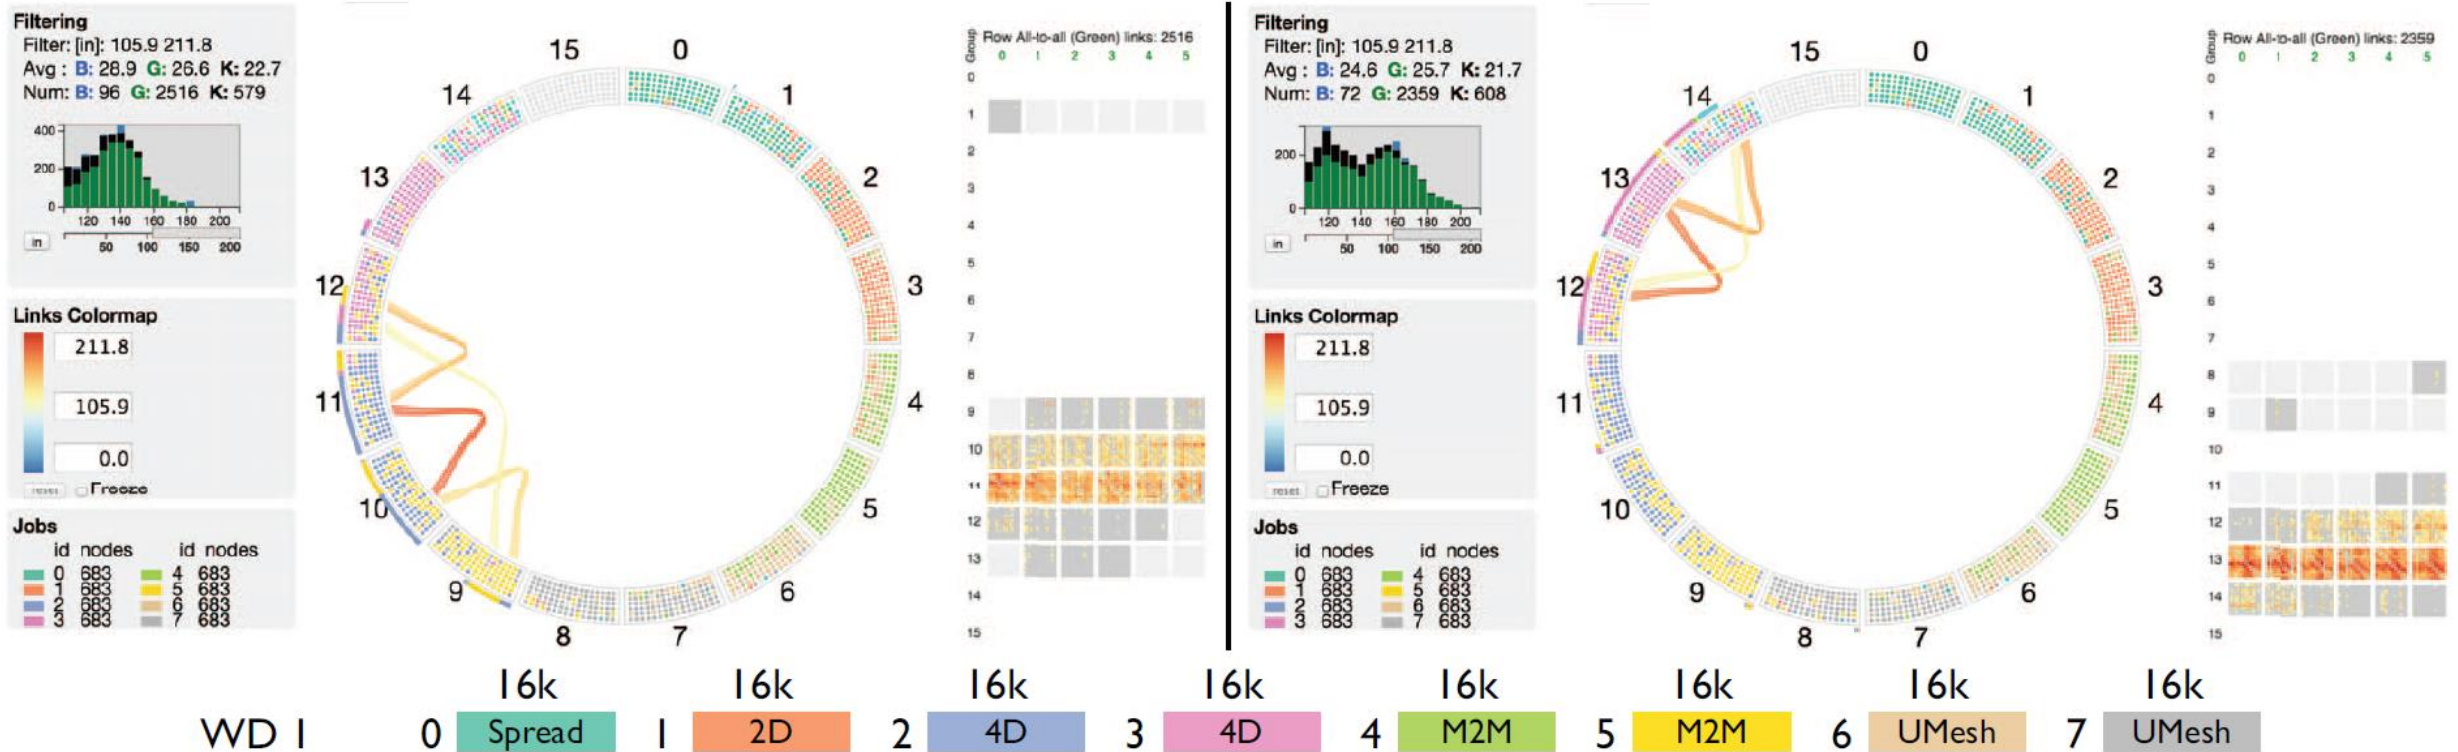
In scattered

- More blue links used (1589)
- Lower maximum traffic (191 MB)

In compact

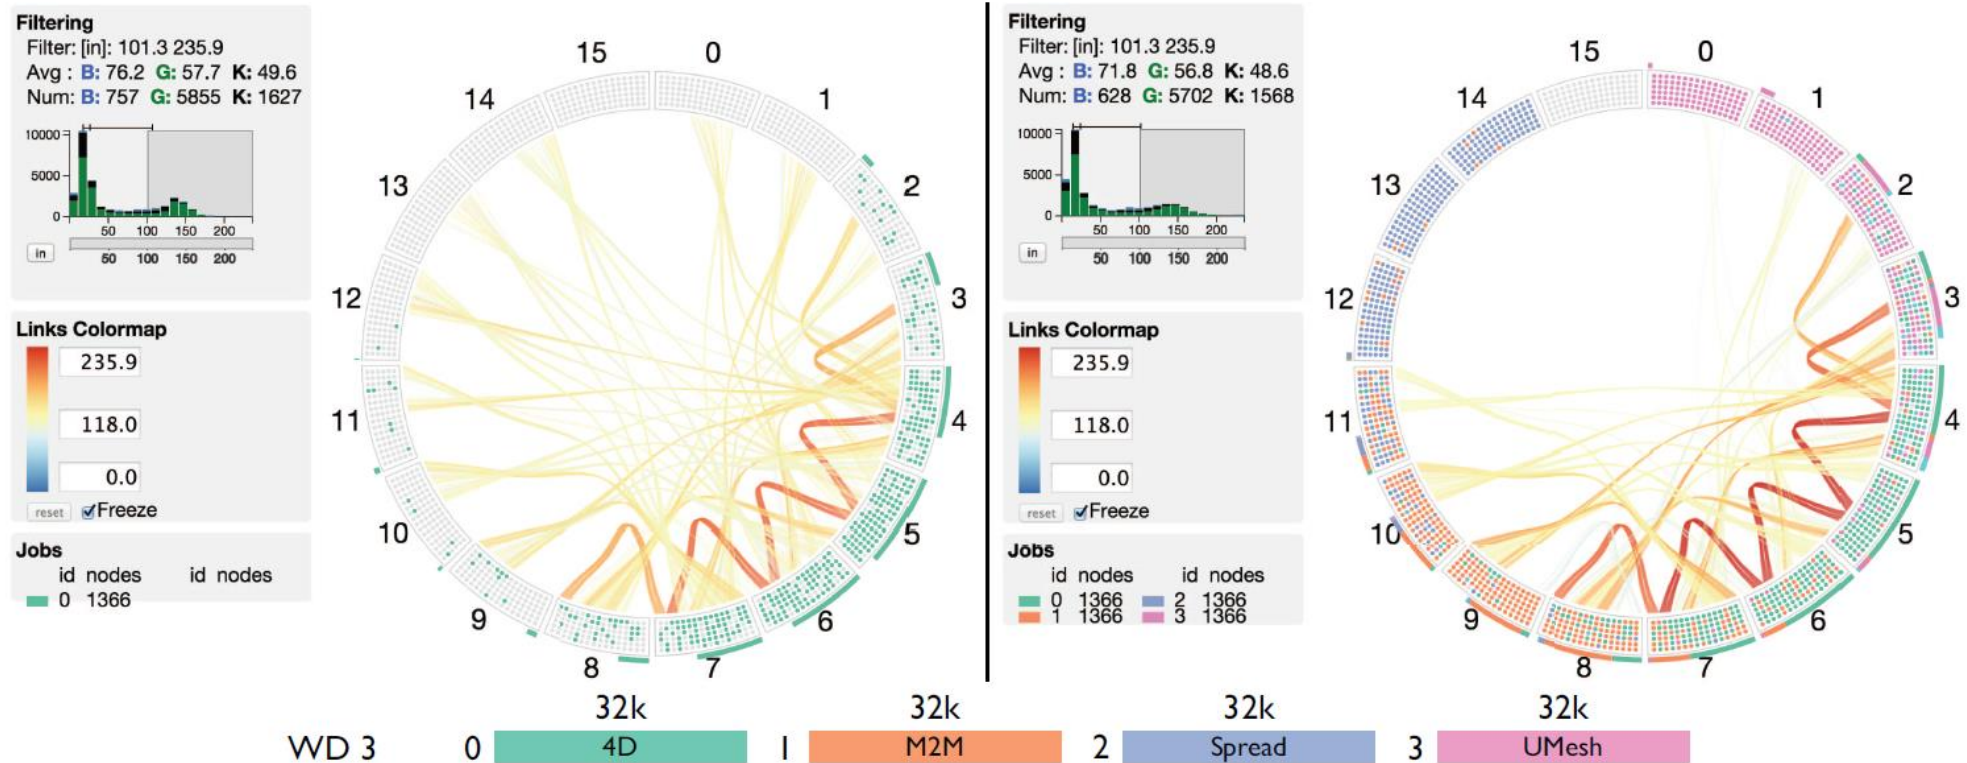
- Less blue links used (1226)
- Higher maximum traffic (236 MB)

# Job Placement Effects (Relative)



Scattered, Job 2 (left) and compact, Job 3 (right) job placements (4D stencil on 16K cores with other jobs). Links shared with communication-heavy code (Job 2 and 3), 8% higher maximum traffic (right figure).

# Inter-job Interference



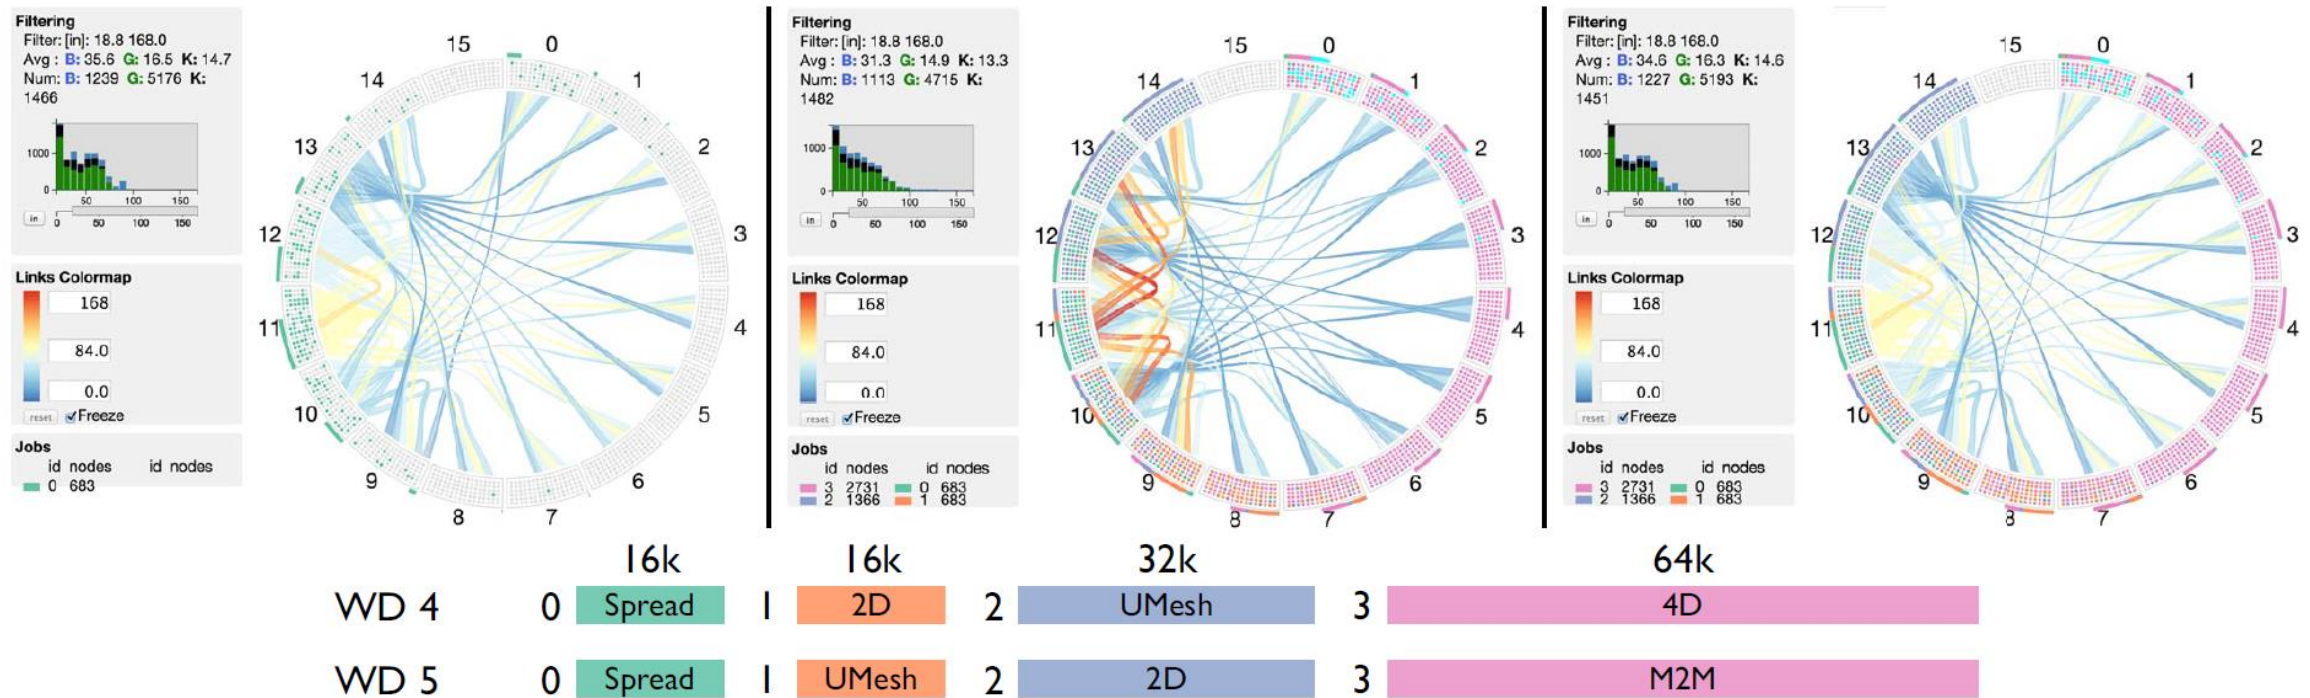
4D stencil (job 0) runs in isolation (left) and 4D stencil (job 0) runs with other jobs (right)

More blue links utilized and lower maximum traffic when other jobs are not running (left)

Job 0's traffic confined to fewer blue links to share bandwidth with other jobs (right)



# Inter-job Interference (Dependent on Job Type)



Congestion increases more when Spread is run with communication-intensive 4D stencil (middle) than when run along with M2M (right)