Profiling - II

Lecture 24

April 15, 2024

Gprof

"gprof: A Call Graph Execution Profiler", by S. Graham, P. Kessler, M. McKusick; Proceedings of the SIGPLAN '82 Symposium on Compiler Construction, SIGPLAN Notices, Vol. 17, No 6, pp. 120-126, June 1982.

- Compile with –g –pg flags
- gprof ./exe gmon.out > gprof.out

Sections in the mpiP report

- 31:@--- MPI Time (seconds) ----------
- 52:@--- Callsites: 43 ----------
- 99:@--- Aggregate Time (top twenty, descending, milliseconds) -------
- 123:@--- Aggregate Sent Message Size (top twenty, descending, bytes) ---------
- 146:@--- Callsite Time statistics (all, milliseconds): 688 ----------
- 923:@--- Callsite Message Sent statistics (all, sent bytes) ----------
- 1268:@--- End of Report ----------

Aggregate Time – Strong Scaling Comparison (mg)

@--- Aggregate Time (top twenty, descending, milliseconds)

Call	Site	Time	App%	MPI%	COV
Bcast	12	22.5	0.49	19.71	0.66
Send	9	21.1	0.46	18.47	0.10
Send	20	14.1	0.31	12.39	0.02
Send	1	13.4	0.29	11.74	0.32

Processes = 4

@--- Aggregate Time (top twenty, descending, milliseconds)

Call	Site	Time	App%	MPI%	COV
Barrier	7	149	2.24	21.84	0.01
Send	9	140	2.10	20.48	0.81
Send	21	123	1.84	17.94	0.87
Wait	26	58.8	0.88	8.60	0.09

Processes = 16

Aggregate Time – Strong Scaling Comparison (ft)

@--- Aggregate Time (top twenty, descending, milliseconds)

Call	Site	Time	App%	MPI%	COV
Alltoall	9	443	5.69	84.89	0.03
Bcast	7	43.3	0.56	8.29	0.00
Reduce	4	32.4	0.42	6.21	0.02
Barrier	5	1.57	0.02	0.30	1.16

Processes = 4

@--- Aggregate Time (top twenty, descending, milliseconds)

Call	Site	Time	App%	MPI%	COV
Alltoall	9	1.73e+03	16.22	91.43	0.05
Reduce	4	76.4	0.72	4.03	0.84
Comm_split	10	44.3	0.41	2.34	0.92
Bcast	7	24.8	0.23	1.31	0.48

Processes = 16

Aggregate Time – Data Scaling (cg on 16 processes)

@--- Aggregate Time (top twenty, descending, milliseconds)

Call	Site	Time	App%	MPI%	COV
Bcast	3	477	16.58	39.24	0.01
Wait	21	176	6.11	14.46	0.74
Send	10	162	5.64	13.34	0.85
Wait	6	89.6	3.11	7.37	0.13

Class = A (small problem)

@--- Aggregate Time (top twenty, descending, milliseconds)

Call	Site	Time	App%	MPI%	COV
Wait	21	1.03e+04	3.07	31.48	0.79
Send	15	8.84e+03	2.65	27.13	0.19
Send	12	8.44e+03	2.53	25.91	0.79
Wait	11	728	0.22	2.23	1.49

Class = C (large problem)

MPI Time vs. App Time (Class = A NPROCS=16)

@--- Aggregate Time (top twenty, descending, milliseconds)

Call	Site	Time	App%	MPI%	COV
Bcast	3	477	16.58	39.24	0.01
Allreduce	5	246	1.31	77.45	0.60
Alltoall	9	1.73e+03	16.22	91.43	0.05
Recv	3	3.78e+03	5.30	27.16	0.76
Barrier	7	149	2.24	21.84	0.01
Waitall	41	1.89e+03	2.13	20.52	0.17

cg
ep
ft
lu
mg
sp

IPM Profiles

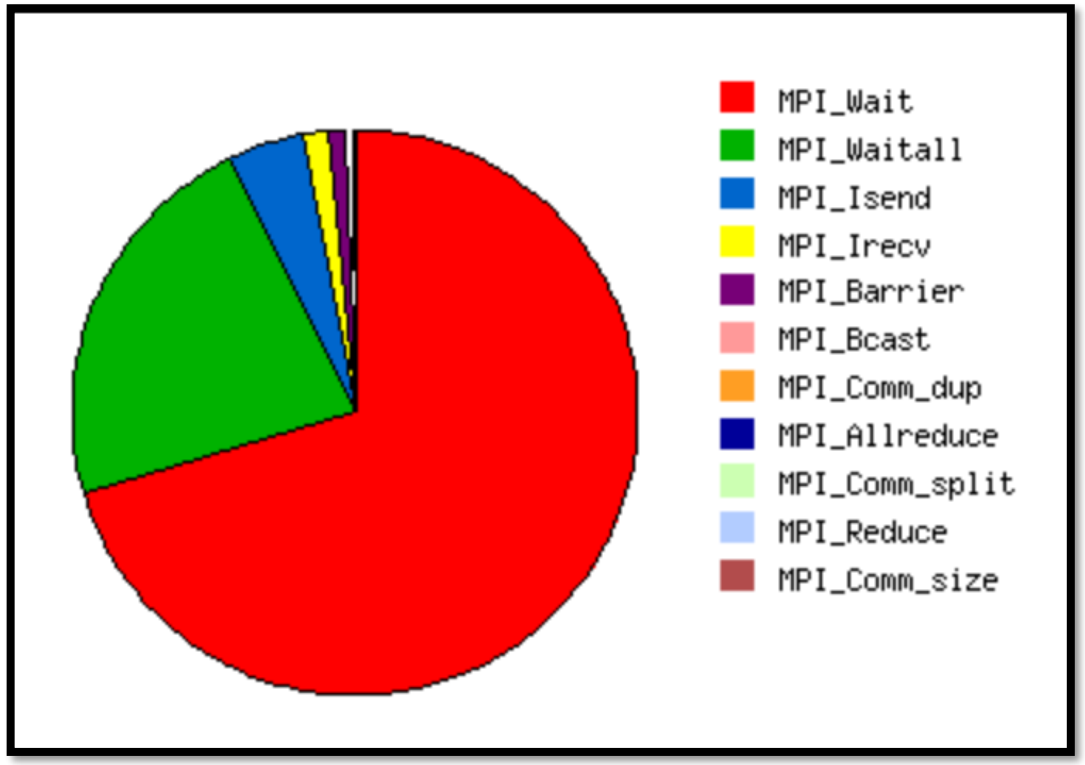

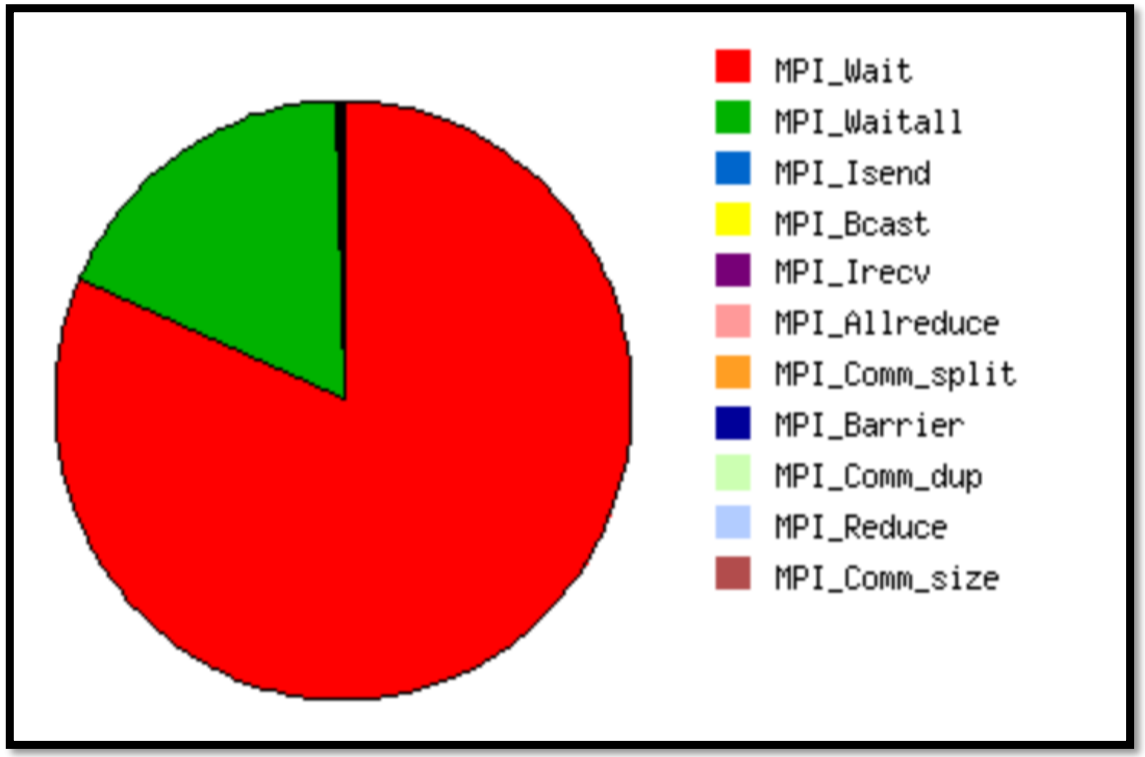

Configuration difference?

IMB Reduce (NPROCS = 4)

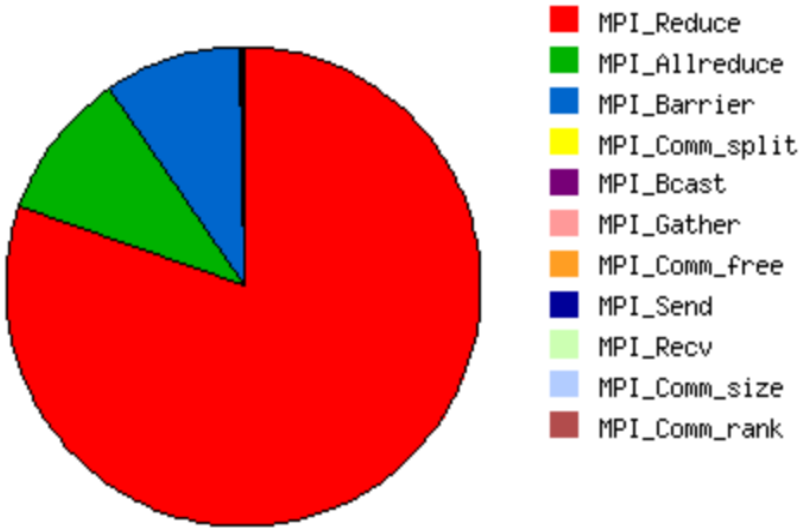

Communication Event Statistics (100.00% detail, 7.4219e-07 error)			
	Buffer Size	**Ncalls**	**Total Time**
MPI_Reduce	32768	12006	0.222
MPI_Allreduce	4	308	0.217
MPI_Barrier	0	984	0.205
MPI_Reduce	4194304	126	0.203
MPI_Reduce	2097152	246	0.181
MPI_Reduce	65536	7686	0.167
MPI_Reduce	131072	3846	0.154
MPI_Reduce	1048576	486	0.147
MPI_Reduce	262144	1926	0.145
MPI_Reduce	524288	966	0.141
MPI_Reduce	16384	12006	0.121
MPI_Reduce	8192	12006	0.065
MPI_Reduce	4096	12006	0.039
MPI_Reduce	2048	12006	0.028
MPI_Reduce	1024	12006	0.023
MPI_Reduce	512	12006	0.019

IMB Reduce (NPROCS = 4)

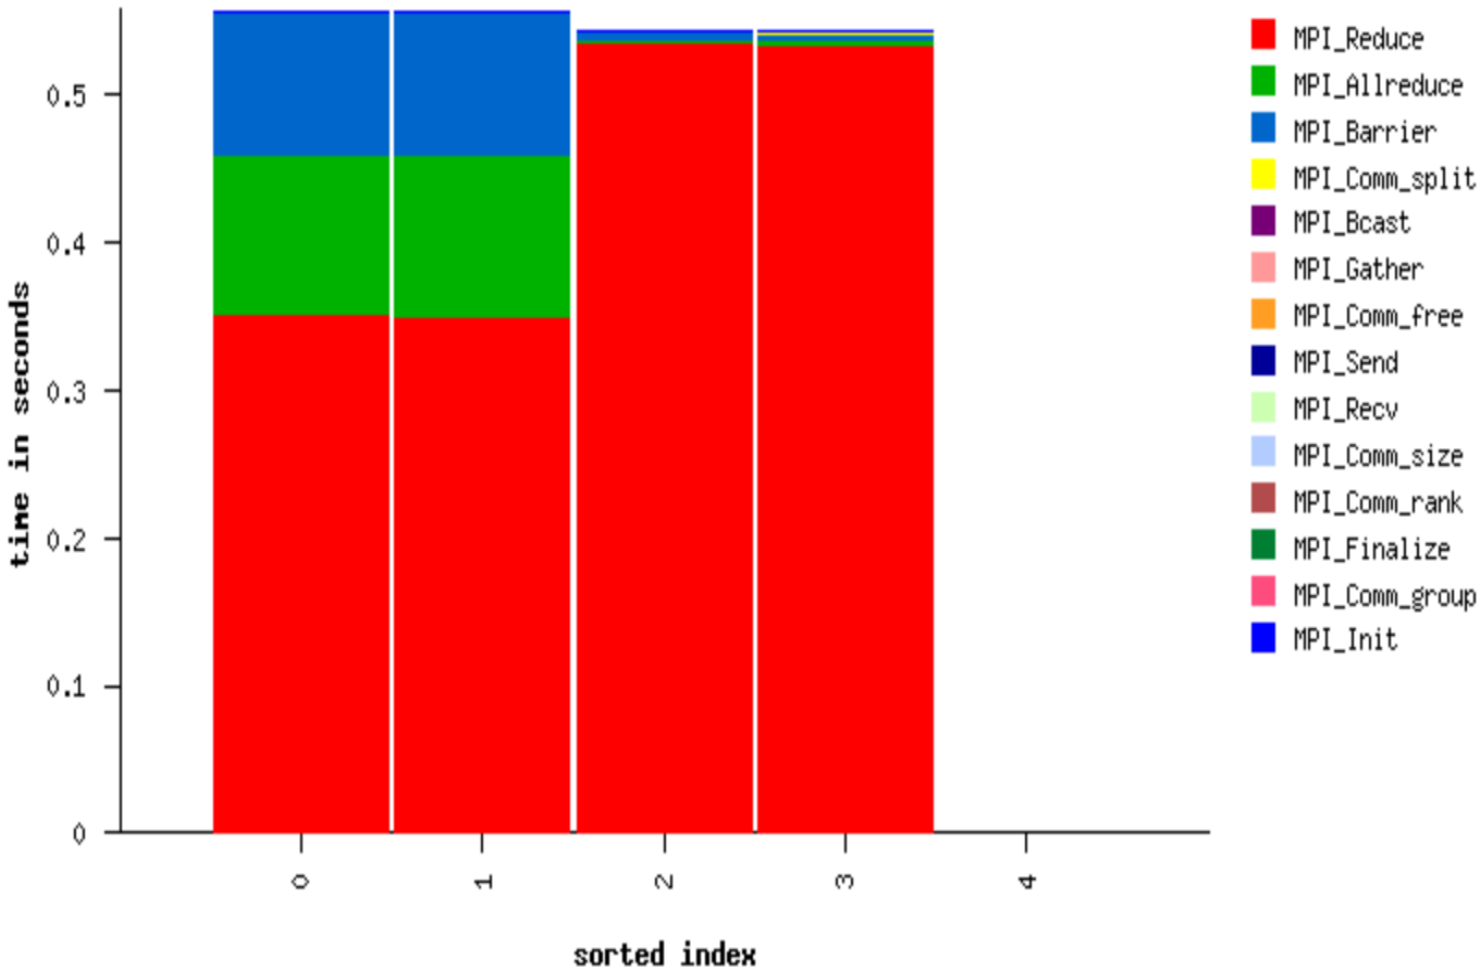

Communication Matrix (IMB Reduce, NPROCS=4)

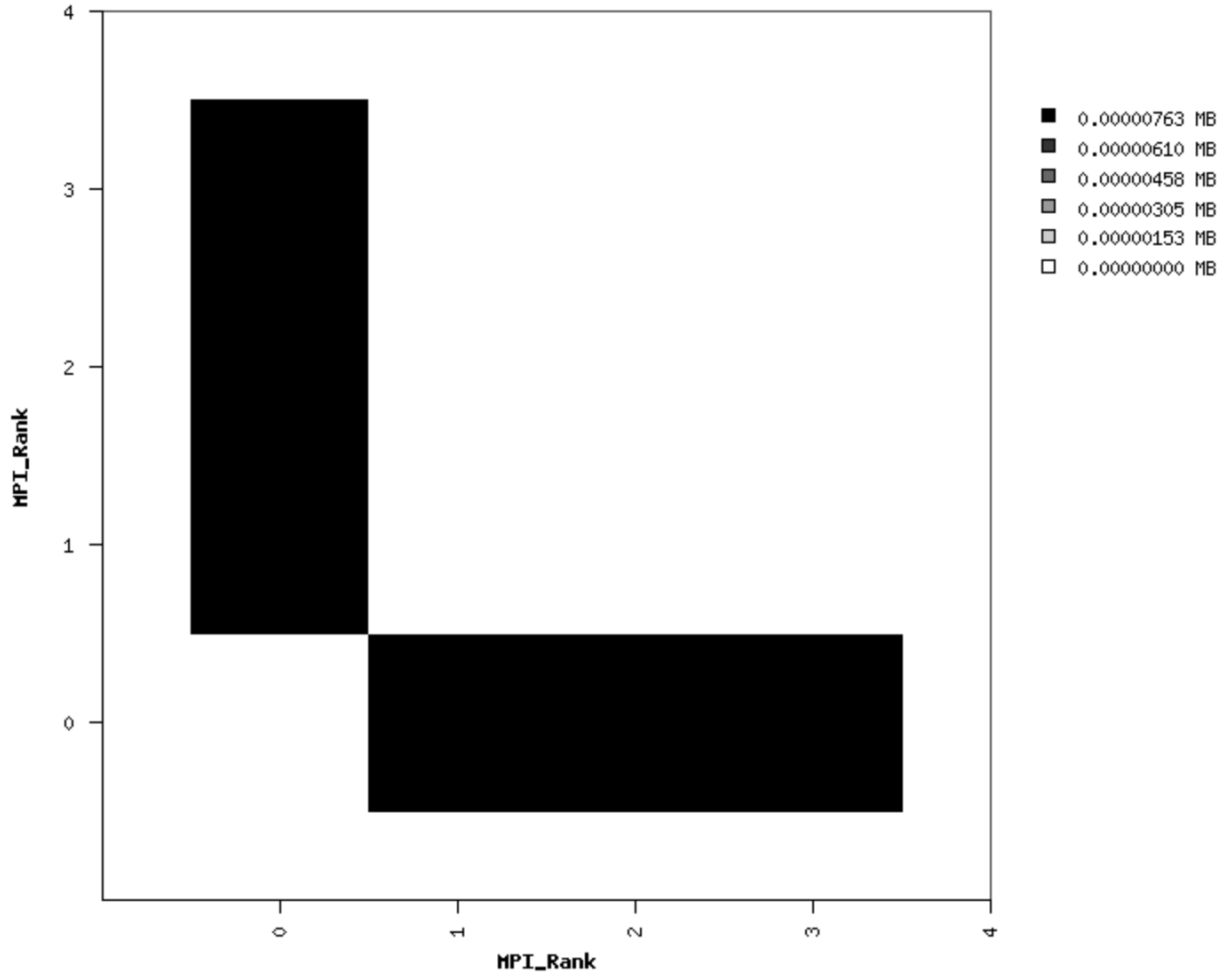

IMB Reduce (NPROCS = 8, 1 host)

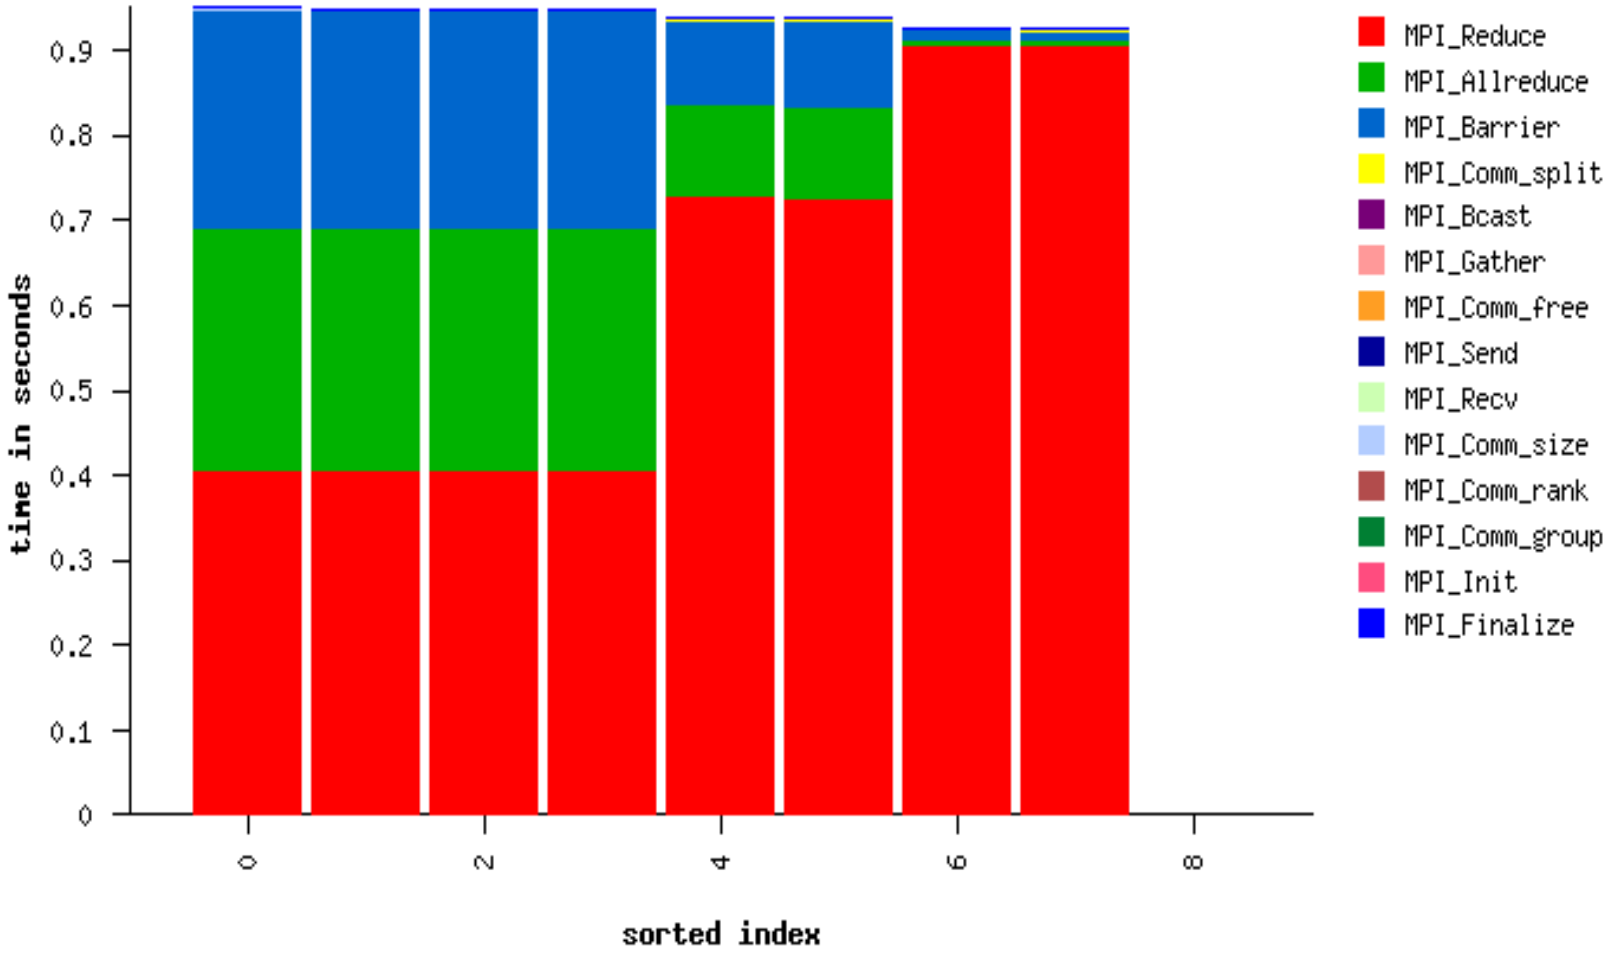

IMB Reduce (NPROCS = 16, 2 hosts)

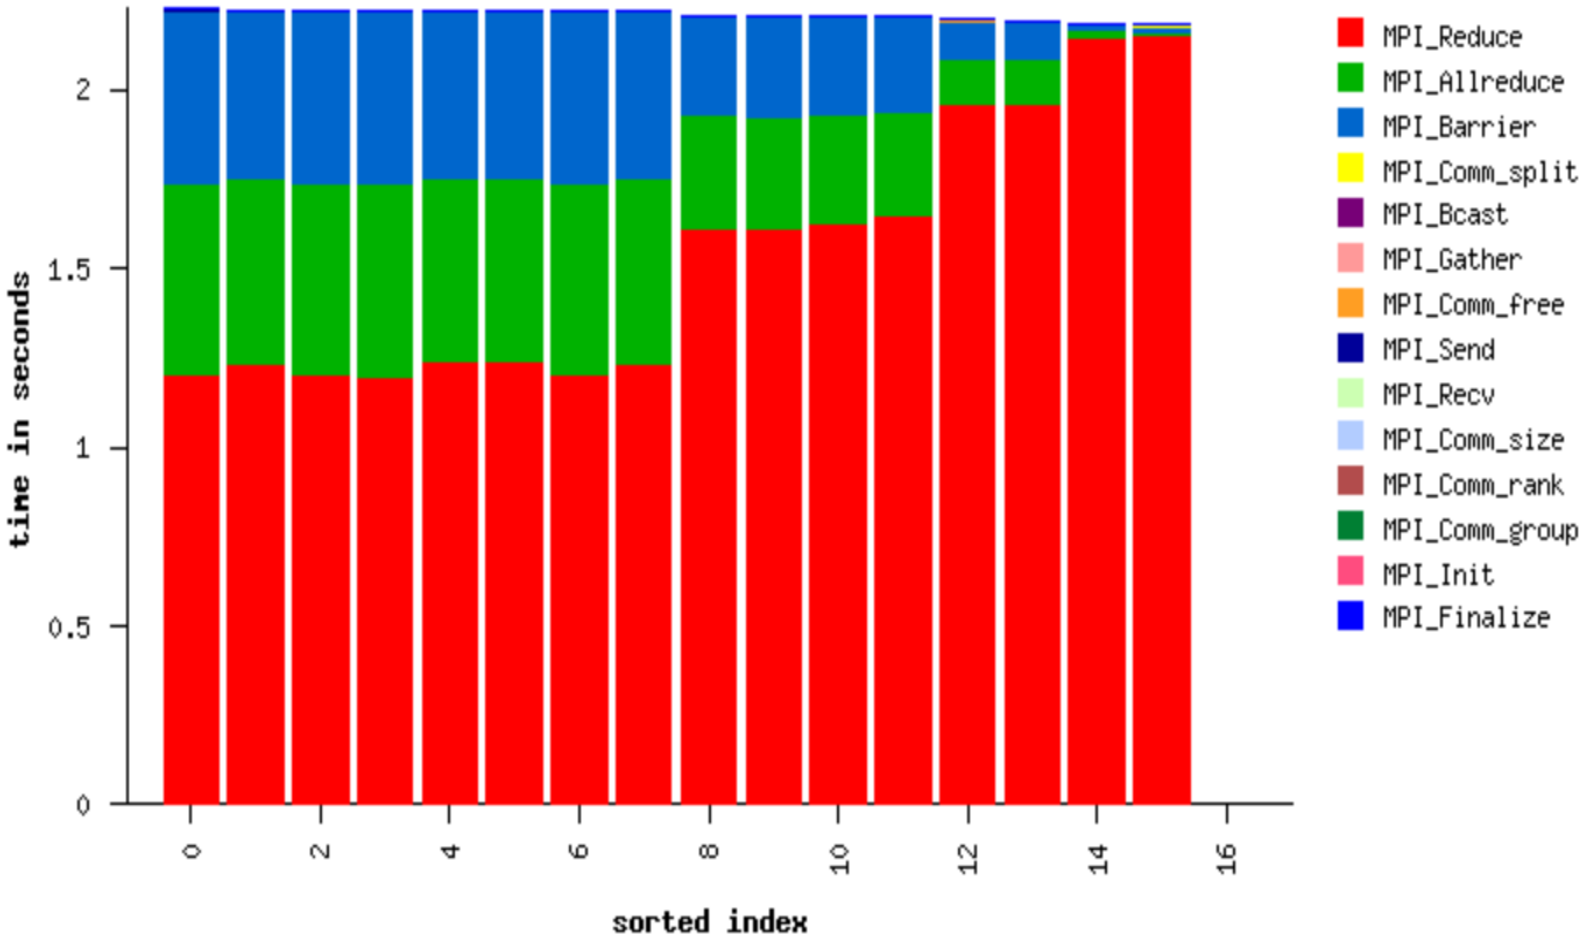

IMB Gather (NPROCS=32, 4 hosts)

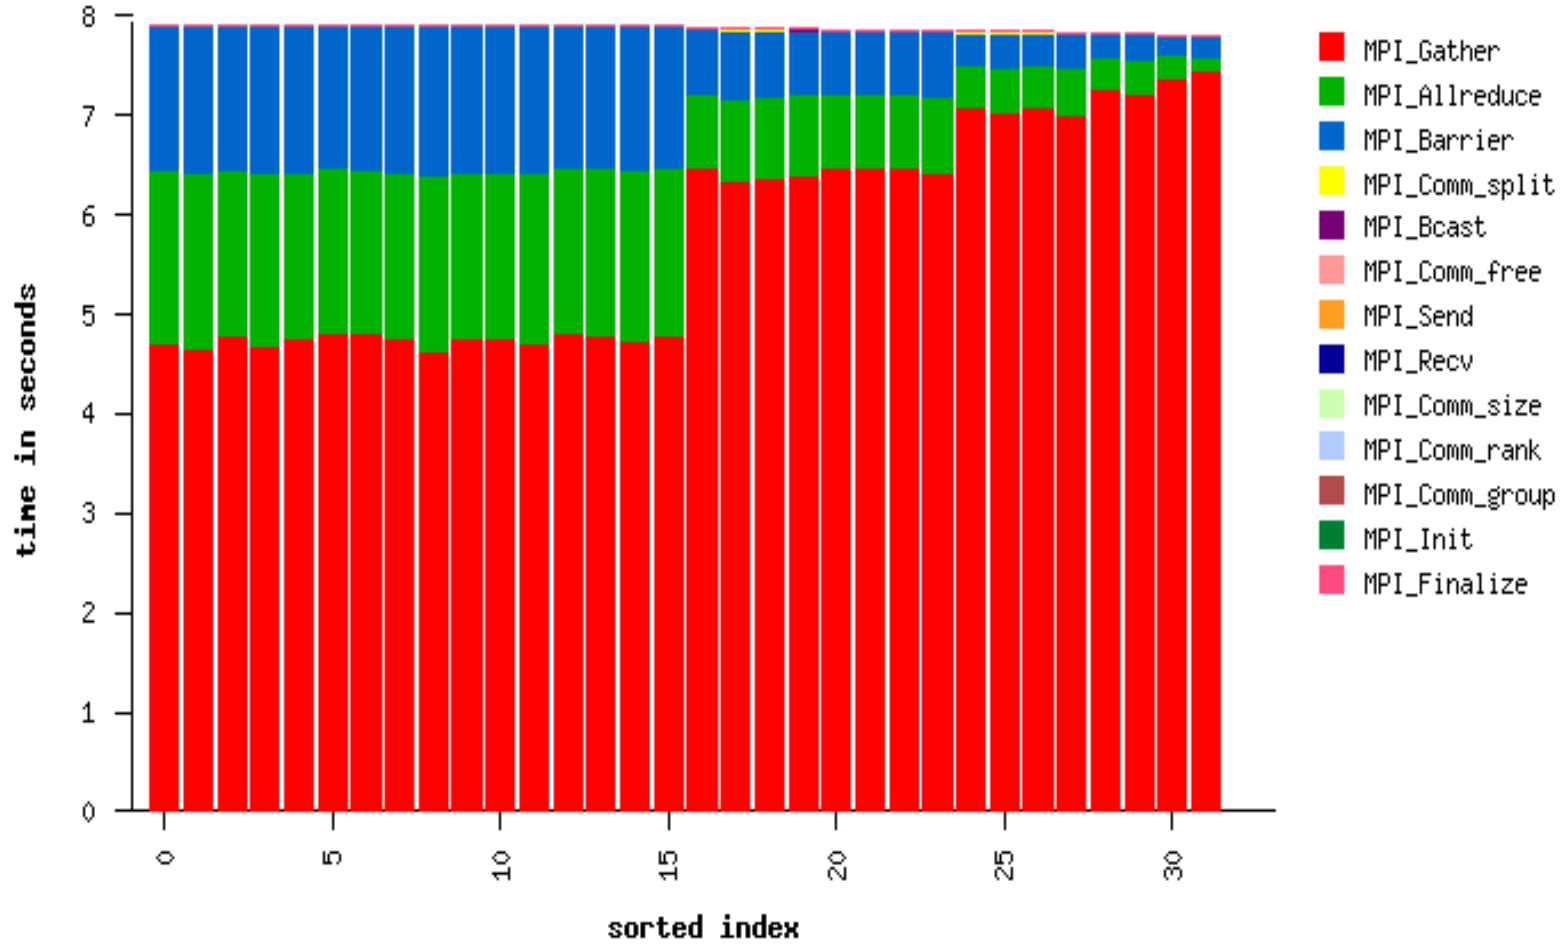

Darshan Internals

- Intercepts MPI-IO routines using PMPI interface
- Data recorded on each process at run time and then merged and stored during MPI_Finalize
- MPI_Wtime() collects timing information
- In-memory file record
  - Array of counters for I/O calls
  - Frequency count of common access sizes

- Dynamic linking at runtime
  - LD_PRELOAD – enables overriding
- Static linking at compile-time
  - Inserting wrapper functions
  - --wrap option
- Time Overhead
  - MPI_Wtime() call - 165 ns
  - Function wrapping - 14 ns *
- Memory overhead
  - File record 2 MB limit per process
  - Aggregate statistics beyond limit

* “24/7 Characterization of Petascale I/O Workloads”

Darshan I/O Profiler

- cd io
- export DARSHAN_LOGPATH=darshan-logs
- mpiicc –o indepIO indepIO.c
- export LD_PRELOAD=../lib (path to libdarshan.so)
- qsub subindepIO.c
- mkdir $DARSHAN_LOGPATH/2024/4/15
- ls –t $DARSHAN_LOGPATH/2024/4/15 [Look for .darshan]
- ./darshan-parser <logfile> > parsed
- grep POSIX_F_FASTEST_RANK_TIME parsed
- grep MPIIO_F_FASTEST_RANK_TIME parsed
- grep MPIIO_F_SLOWEST_RANK_TIME parsed

Revision Q1

MPI_Bcast of 10 KB data (root=0) on the 2D mesh.
There are 8 processes placed on the 4 nodes.
Ranks 0 and 1 are placed on node 1, ranks 2 and 3 are placed on node 2 and so on.
Bandwidth of every link is 1 Gbps. Assume hop=0 between processes in a node. Assume XY routing policy (i.e. messages first traverse in x-dimension, followed by y-dimension).
Total time = 4 ms
Analyze and discuss the effective bandwidth, maximum #hops, and link contention with Bcast.

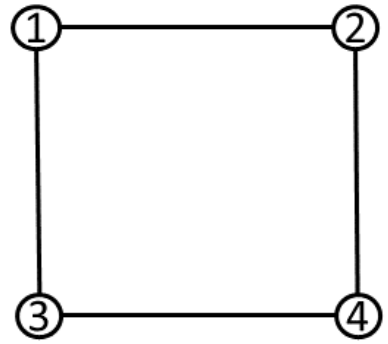

Revision Q2

Compare and contrast recursive doubling algorithm for MPI_Reduce on 8 processes for the following node allocations:

(a) Ranks 0 – 3 are on csews1, ranks 4 – 7 are on csews2

(b) Even ranks are on csews1, odd ranks are on csews2

Revision Q3: 3D domain decomposition

```
17  //initialize
18  for (int i=0; i<N; i++)
19   for (int j=0; j<N; j++)
20    for (int k=0; k<N; k++)
21     data[i][j][k] = (rank+1) * (i+j+k);

22  int xStart=_____,
yStart=_____,
zStart=_____;

23  int xEnd=_____,
yEnd=_____,
zEnd=_____;
```