

# Profiling – III and Revision



Lecture 25

April 17, 2024

# Profiling

```
for (int i=0; i<50; i++)  
{  
    MPI_Barrier (MPI_COMM_WORLD);  
    MPI_Alltoall(message, arrSize, MPI_INT, recvMessage, arrSize, MPI_INT,  
MPI_COMM_WORLD);  
}
```

# Alltoall - NPROCS=4, Data size = 4 KB

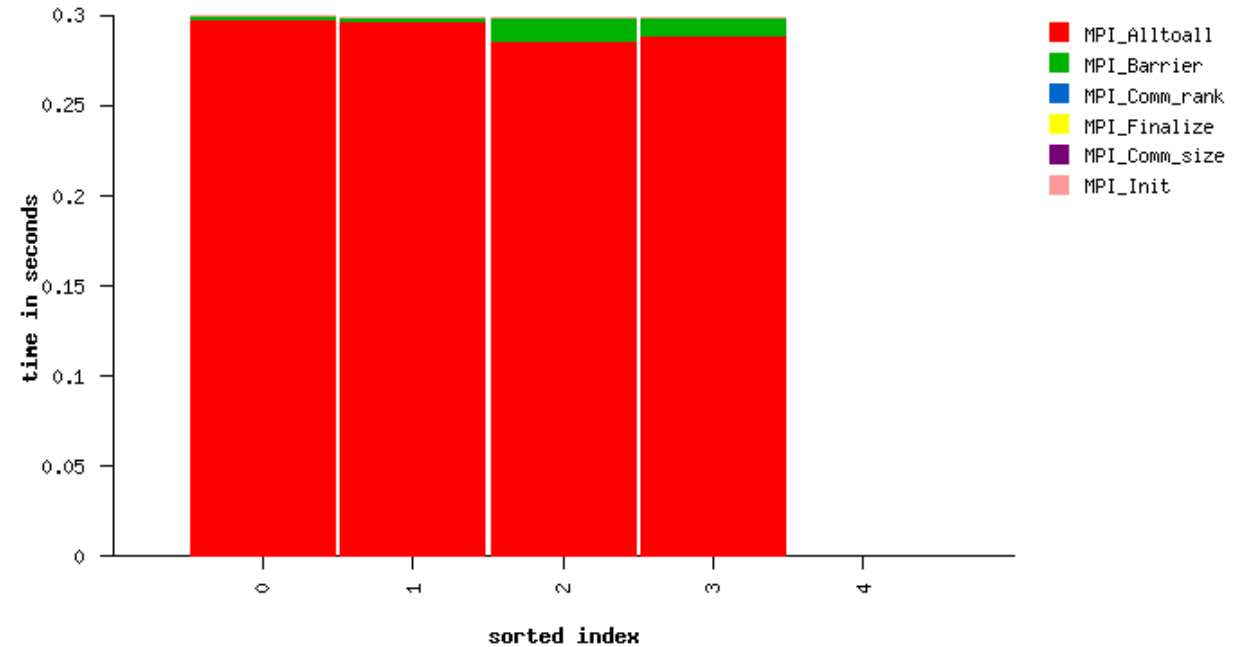


Max. barrier time: 0.2 ms



# Alltoall - NPROCS=4, Data size = 4 MB

Max. barrier time: 1.2 ms


Max. Alltoall time: 7.8 ms



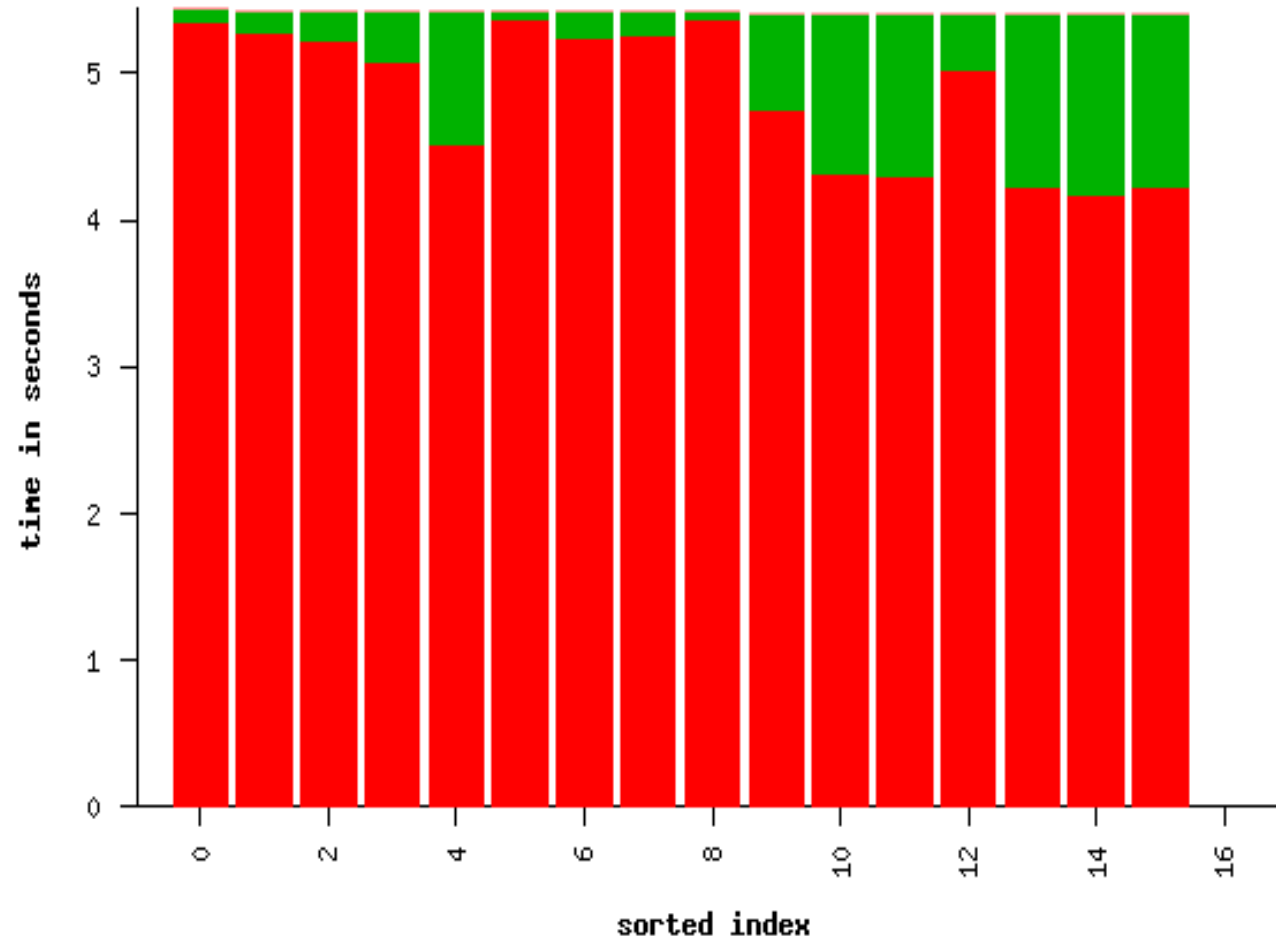
# Alltoall - NPROCS=16, Data size = 4 KB



# Alltoall - NPROCS=16, Data size = 4 MB




■ MPI\_Alltoall  
■ MPI\_Barrier  
■ MPI\_Comm\_rank  
■ MPI\_Comm\_size




■ MPI\_Alltoall  
■ MPI\_Barrier  
■ MPI\_Comm\_rank  
■ MPI\_Comm\_size  
■ MPI\_Finalize  
■ MPI\_Init

# Bcast – 16P, 4 KB




# Bcast – 16P, 4 MB

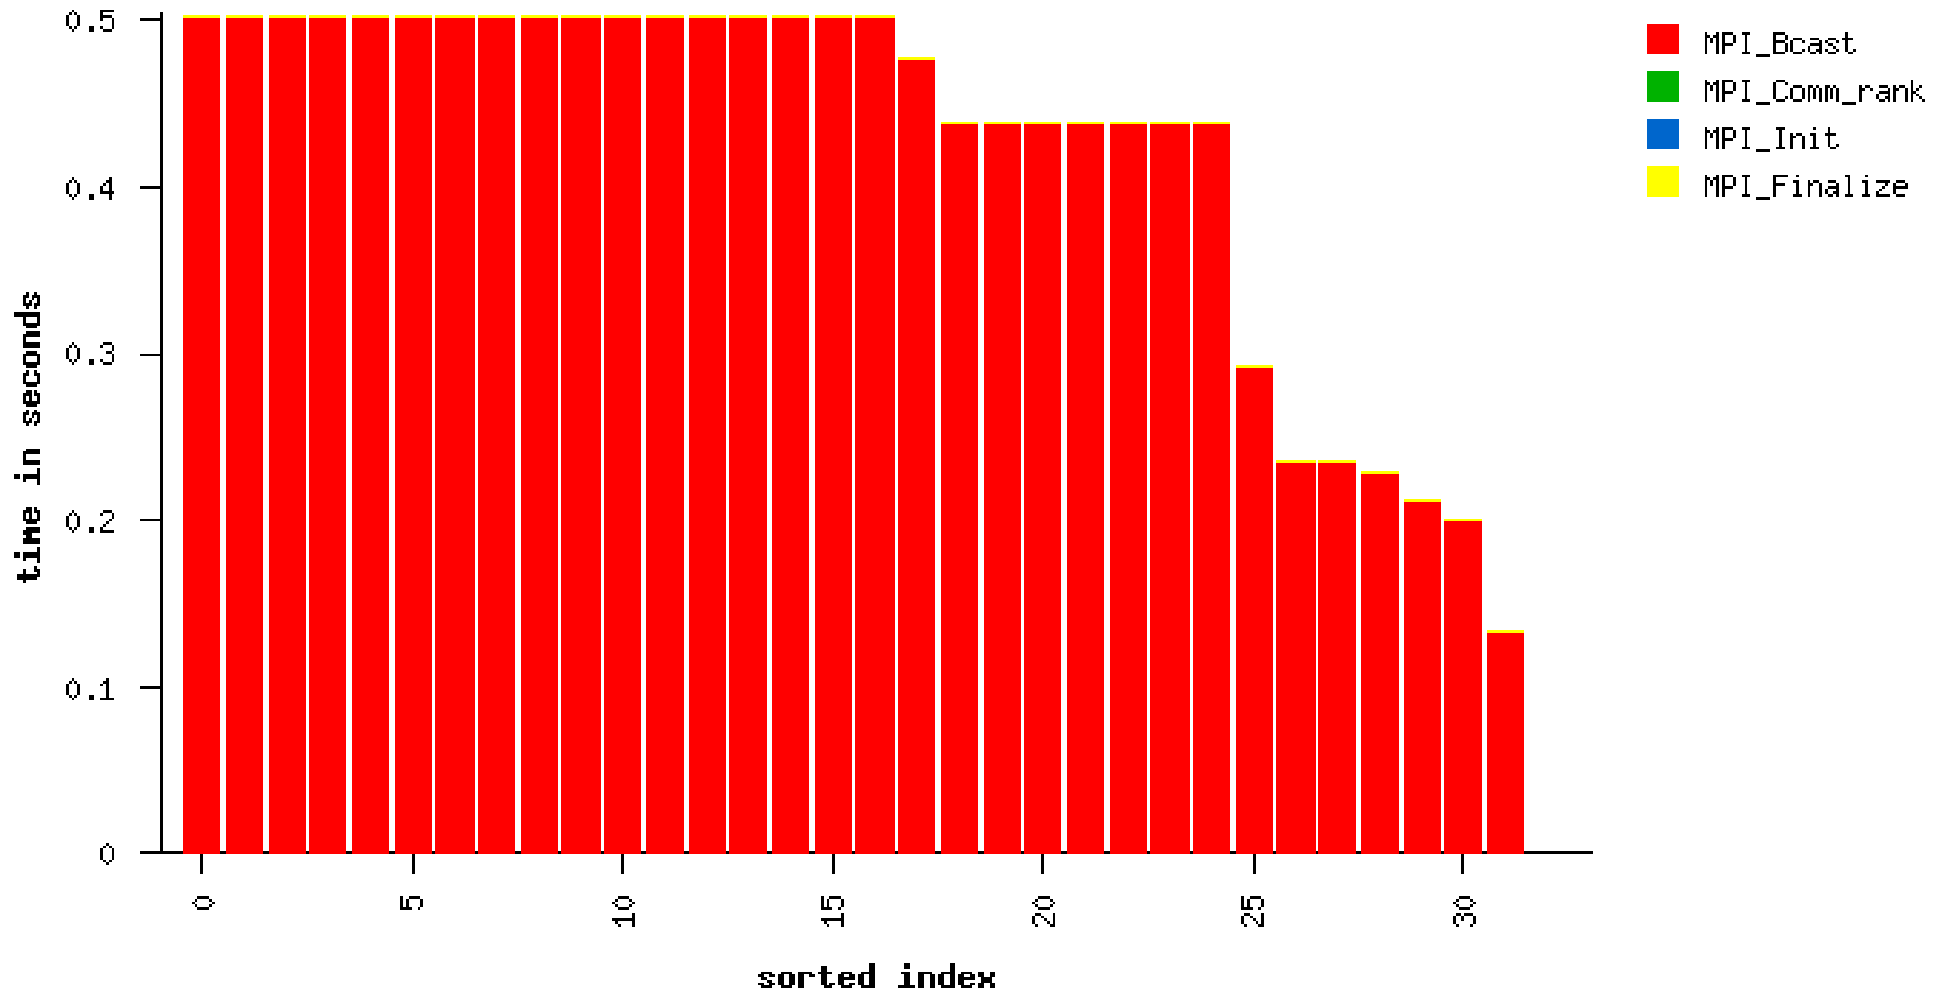




# Bcast – 32P, 4 KB




# Bcast – 32P, 4 MB




# Revision

# Communication Graph Mapping



	512		256		
512				64	256
			512		64
256		512		512	
	64		512		64
	256	64		64	

Linear mapping



Q1: What are the communicating pairs?

Q2: Distance/hops between the communicating pairs?



Q3: Total hop-bytes?

# Estimation Function

- $f_{est}(t, p, M)$  = Cost of placing a task  $t$  onto processor  $p$  under current task mapping  $M$
- Estimate how critical it is to place a task in the current cycle, select the task with maximum criticality
- $T_k$  is the set of tasks yet to be placed
- $P_k$  is the set of processors that are available

$$\begin{array}{l} T_k \cup \bar{T}_k = \emptyset \\ P_k \cup \bar{P}_k = \emptyset \end{array}$$


# Graph Model for SpMV



- Computation load?
  - Similar for both processors
- Number of communications?
  - 8 as per this graph
  - Actually 6

# Parallel I/O

# Access Pattern




Each process reads a **small** chunk of data from a common file

```
MPI_File_set_view (fh, displacement, etype, filetype, "native", info)
MPI_File_read_all (fh, data, datacount, MPI_INT, status)
```



# Multiple Non-contiguous Accesses



- Every process' local array is non-contiguous in file
- Every process needs to make small I/O requests
- Can these requests be merged?

# Revision Q3: 3D domain decomposition

```
17 //initialize
18 for (int i=0; i<N; i++)
19   for (int j=0; j<N; j++)
20     for (int k=0; k<N; k++)
21       data[i][j][k] = (rank+1) * (i+j+k);
```

```
22 int xStart=_____,
   yStart=_____,
   zStart=_____;
```

```
23 int xEnd=_____,
   yEnd=_____,
   zEnd=_____;
```

# Revision Q4

A 3D matrix of size  $N \times N \times N$  was written to the file in the usual XYZ memory order.  $P$  processes read this 3D matrix from a file using parallel I/O following a 1D domain decomposition along Y-axis. Write an MPI code snippet for this (you may ignore the obvious initializations and finalizations). Assume that  $N$  is divisible by  $P$ .

# Revision Q5

A sequential program  $P$  consists of three parts  $A$ ,  $B$ ,  $C$ . Part  $B$  is not parallelizable. Parts  $A$  and  $C$  are parallelizable. The sequential runtimes are  $S_a$ ,  $S_b$ ,  $S_c$  for the parts  $A$ ,  $B$ ,  $C$  respectively. Derive the speedup of  $P$  on  $N$  processes, where the overhead to parallelize part  $A$  is  $O_a$ , overhead to parallelize part  $C$  is  $O_c$ .

# Revision Q6

